

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/356217103>

# Porting a scientific application to GPU using C++ standard parallelism

**Presentation** · November 2021

DOI: 10.13140/RG.2.2.27117.92647

CITATIONS

2

READS

612

7 authors, including:



**Jonas Latt**  
University of Geneva

114 PUBLICATIONS 3,215 CITATIONS

SEE PROFILE



**Christophe Coreixas**  
BNU-HKBU United International College

58 PUBLICATIONS 1,097 CITATIONS

SEE PROFILE

NVIDIA GTC 21  
November 2021

# Porting a scientific application to GPU using C++ standard parallelism

**Presenter:** *Jonas Latt - Université de Genève (UNIGE), Switzerland*  
**Contributors:** *Christophe Coreixas (UNIGE)*  
*Francesco Marson (UNIGE)*  
*Karthik Thyagarajan (UNIGE)*  
*Jose Pedro de Santana Neto (UNIGE)*  
*Sriharan BS (IITM)*  
*Gonzalo Brito (NVIDIA)*

# Outline

## Part I. Introduction

Everything you need to know to follow this presentation: C++ parallel algorithms, lattice Boltzmann method, the Palabos project.

## Part II. Design: From Object-Oriented to Data-Oriented

For the most part, the code can be executed on GPU as is. But for best performance, the overall architecture must be migrated from object-oriented to data-oriented.

## Part III. In-depth discussion: performance improvements

Some neat tricks to get even better performance with little effort.




Part I

# INTRODUCTION

# The Palabos software library


Fluid-Structure



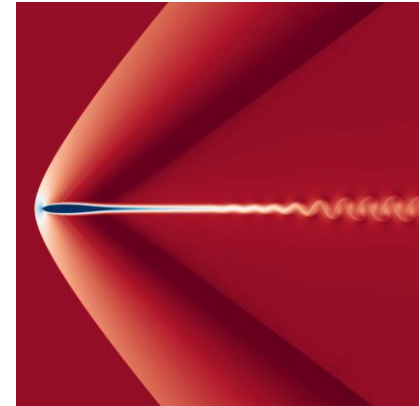
Multi-Phase



Aerodynamics



Supersonic



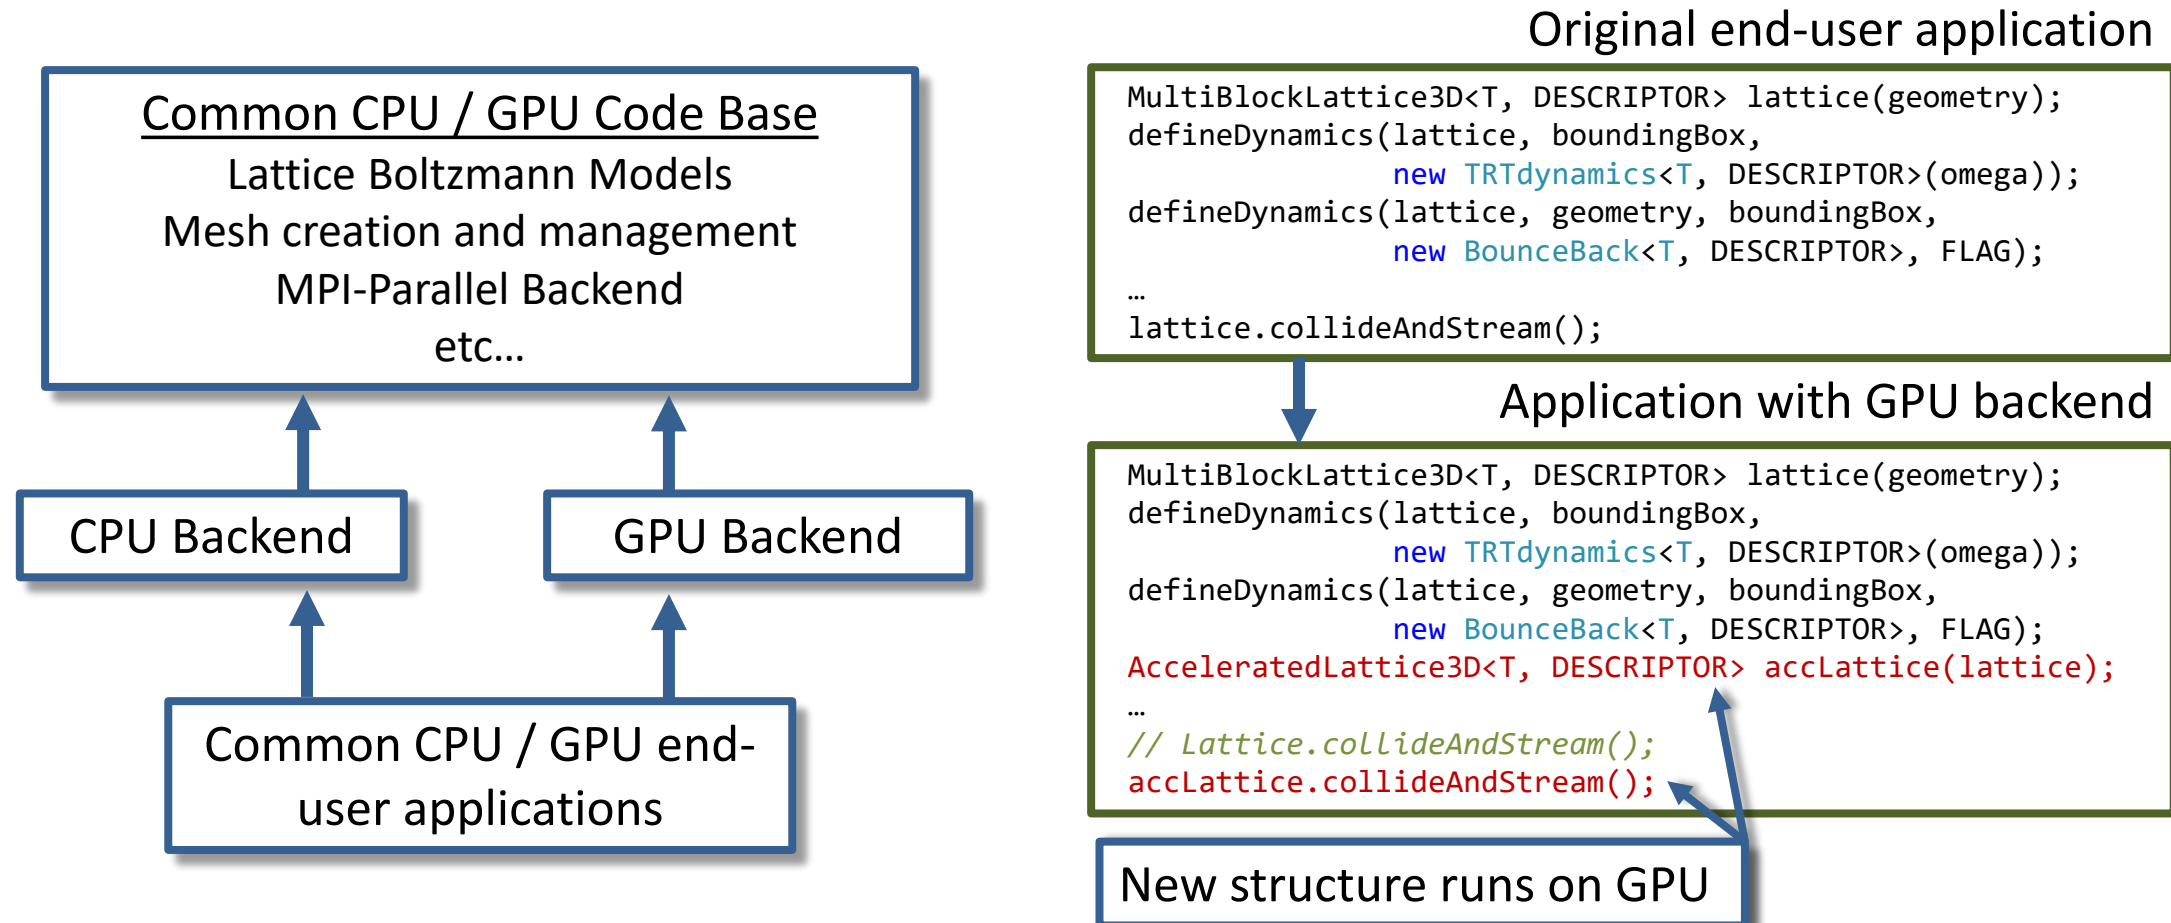
A software library for complex fluid flow simulations

Numerical method: Lattice Boltzmann

Massively parallel (originally CPU)

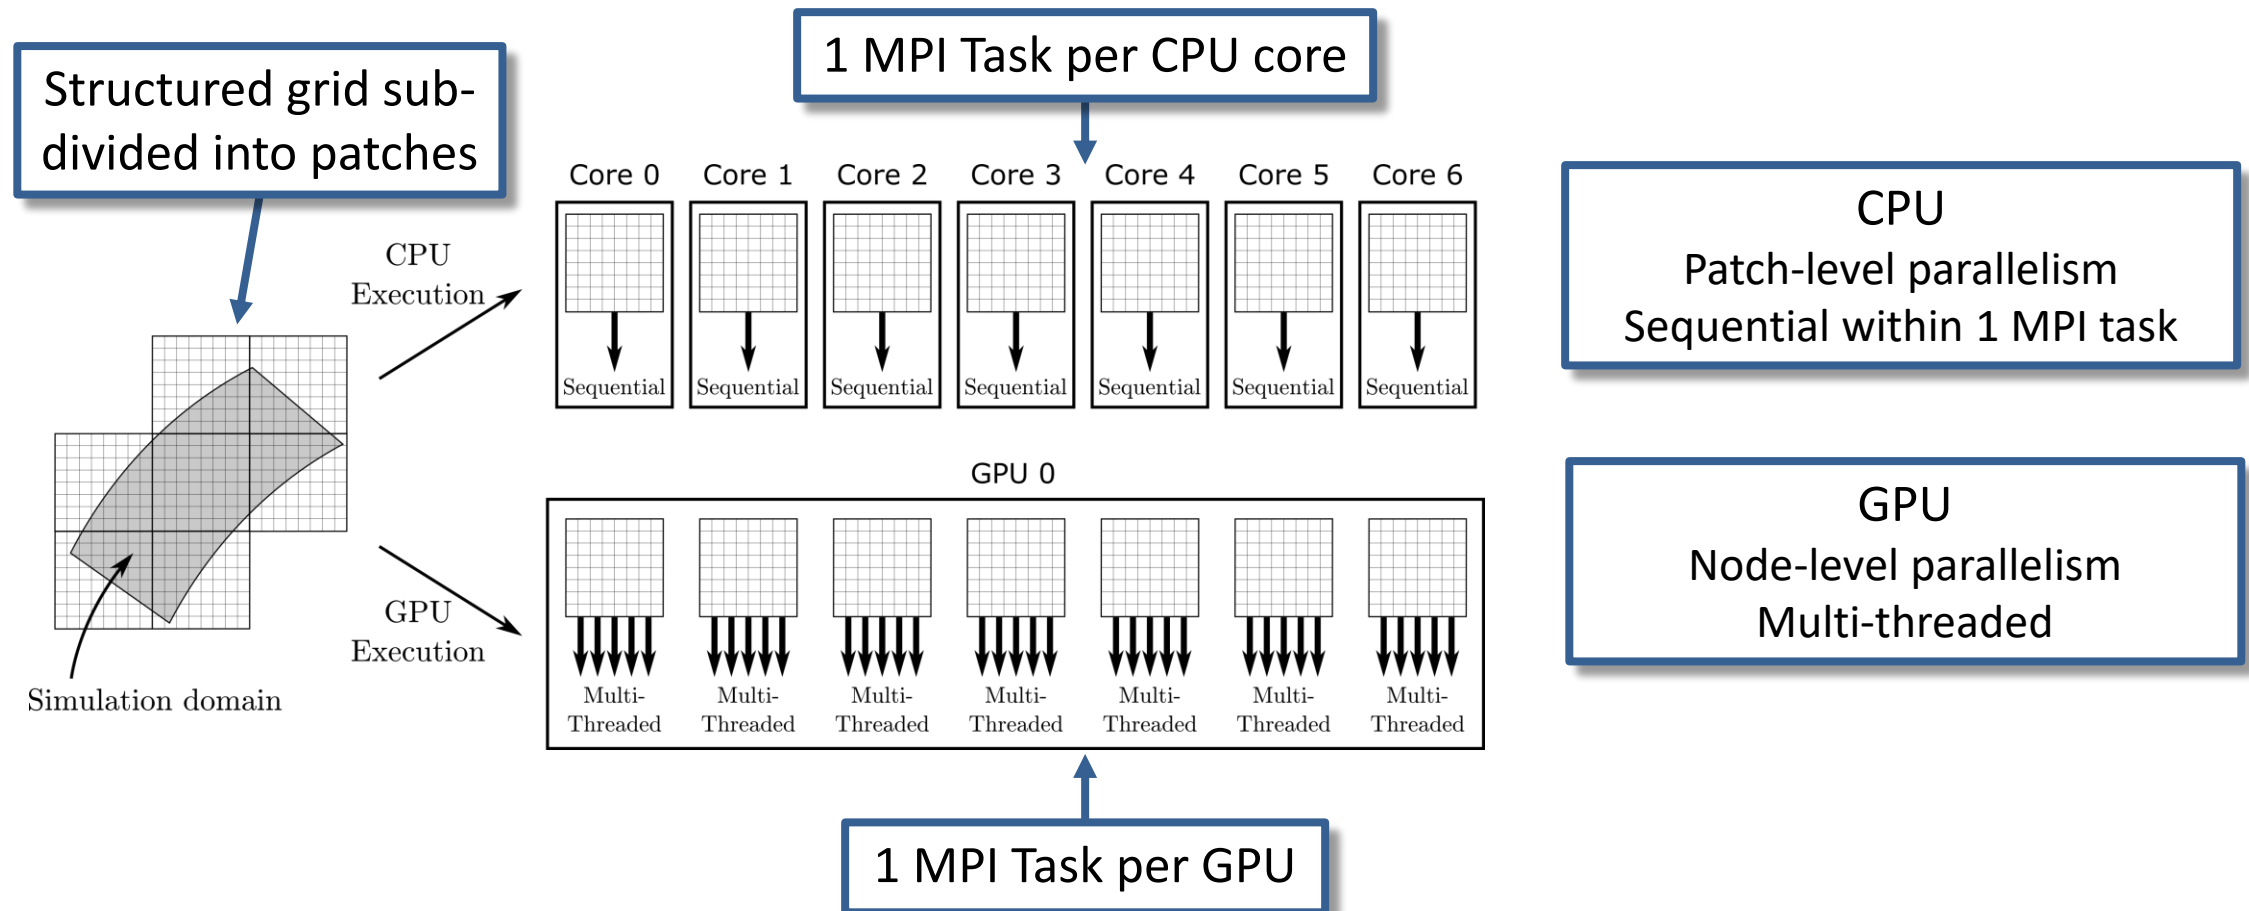
# A GPU backend for Palabos

*Modifications to the existing CPU code base should be minimal*



# A GPU backend for Palabos

*On GPU, multi-threading must be more fine-grained*



# Reminder: C++ Standard Parallelism

*Since C++17, C++ includes parallel language features (no extension, no library needed)*

```
vector<double> v = { 1, 2, 3, 4 }, w(4);  
transform(execution::par_unseq, begin(v), end(v), begin(w),  
          [](double const& element) { return 2. * element; } );  
// w is {2, 4, 6, 8}
```

Execution policy

Read from  $v$ , write into  $w$

Lambda function: defines the element-wise operation applied to  $v$ .

```
nvc++ -stdpar -o program program.cpp
```



# Available implementations

*The code is portable and gets accelerated for multiple types of parallel platforms*

*Common formalism*

C++ Parallel Algorithms

*Hardware-specific  
implementation*

NVIDIA stdpar  
for GPU

Intel Threading  
Building Blocks

GPU  
heterogeneous system

CPU  
homogeneous  
system

# CUDA Unified Memory

*Parallel STL execution on heterogeneous platform is possible thanks to managed memory model*

```
vector<double> v = { 0., 1., 2., 3., 4., 5. };  
for_each(begin(v), end(v), [](double& x) { x = sin(x / N * M_PI); });  
  
for_each( execution::par_unseq, begin(v), end(v),  
          [](double& x) { x = sqrt(x); } );
```

Executed on host

Executed on device

Automatic data transfer

This model encourages hybrid CPU / GPU code and porting a code to GPU progressively.

# Further Resources

## Use of standard language parallelism for GPU programming

GTC21 Spring session:

*Fluid Dynamics on GPUs with C++ Parallel Algorithms*

<https://www.nvidia.com/en-us/on-demand/session/gtcspring21-s32076/>

Current GTC21 November session, on Thursday:

*Accelerated Computing with Standard C++, Python, and Fortran [A31181]*

## Use of C++ parallel algorithms for lattice Boltzmann applications

Open-source code STLBM (reusable code-snippets):

<https://www.gitlab.com/UnigeHPFS/stlbm>

# From CPU to GPU in 80 days

*Case Study: The Palabos fluid solver was ported to GPU in less than three months this summer*

<https://palabos.unige.ch/community/cpu-gpu-80-days>

Project: a GPU-port of the Palabos library  
(around half a million lines of code)

Goal: A GPU backend of Palabos with

- Same ease of maintenance as provided by the CPU backend
- Same ease of use as provided by the CPU backend



# Replace loops by for\_each: is this it?

```
for (int iX = 0; iX < nx; ++iX) {  
    for (int iY = 0; iY < ny; ++iY) {  
        for (int iZ = 0; iZ < nz; ++iZ) {  
            cell(iX, iY, iZ).collideAndStream();  
        }  
    }  
}
```



```
for_each(execution::par_unseq,  
         begin(pop), end(pop), [](Cell& cell)  
{  
    cell.collideAndStream();  
} );
```

Is this it? Code runs  
on GPU after simple  
substitution ?

Not entirely...

This presentation  
highlights some  
points to keep in  
mind.

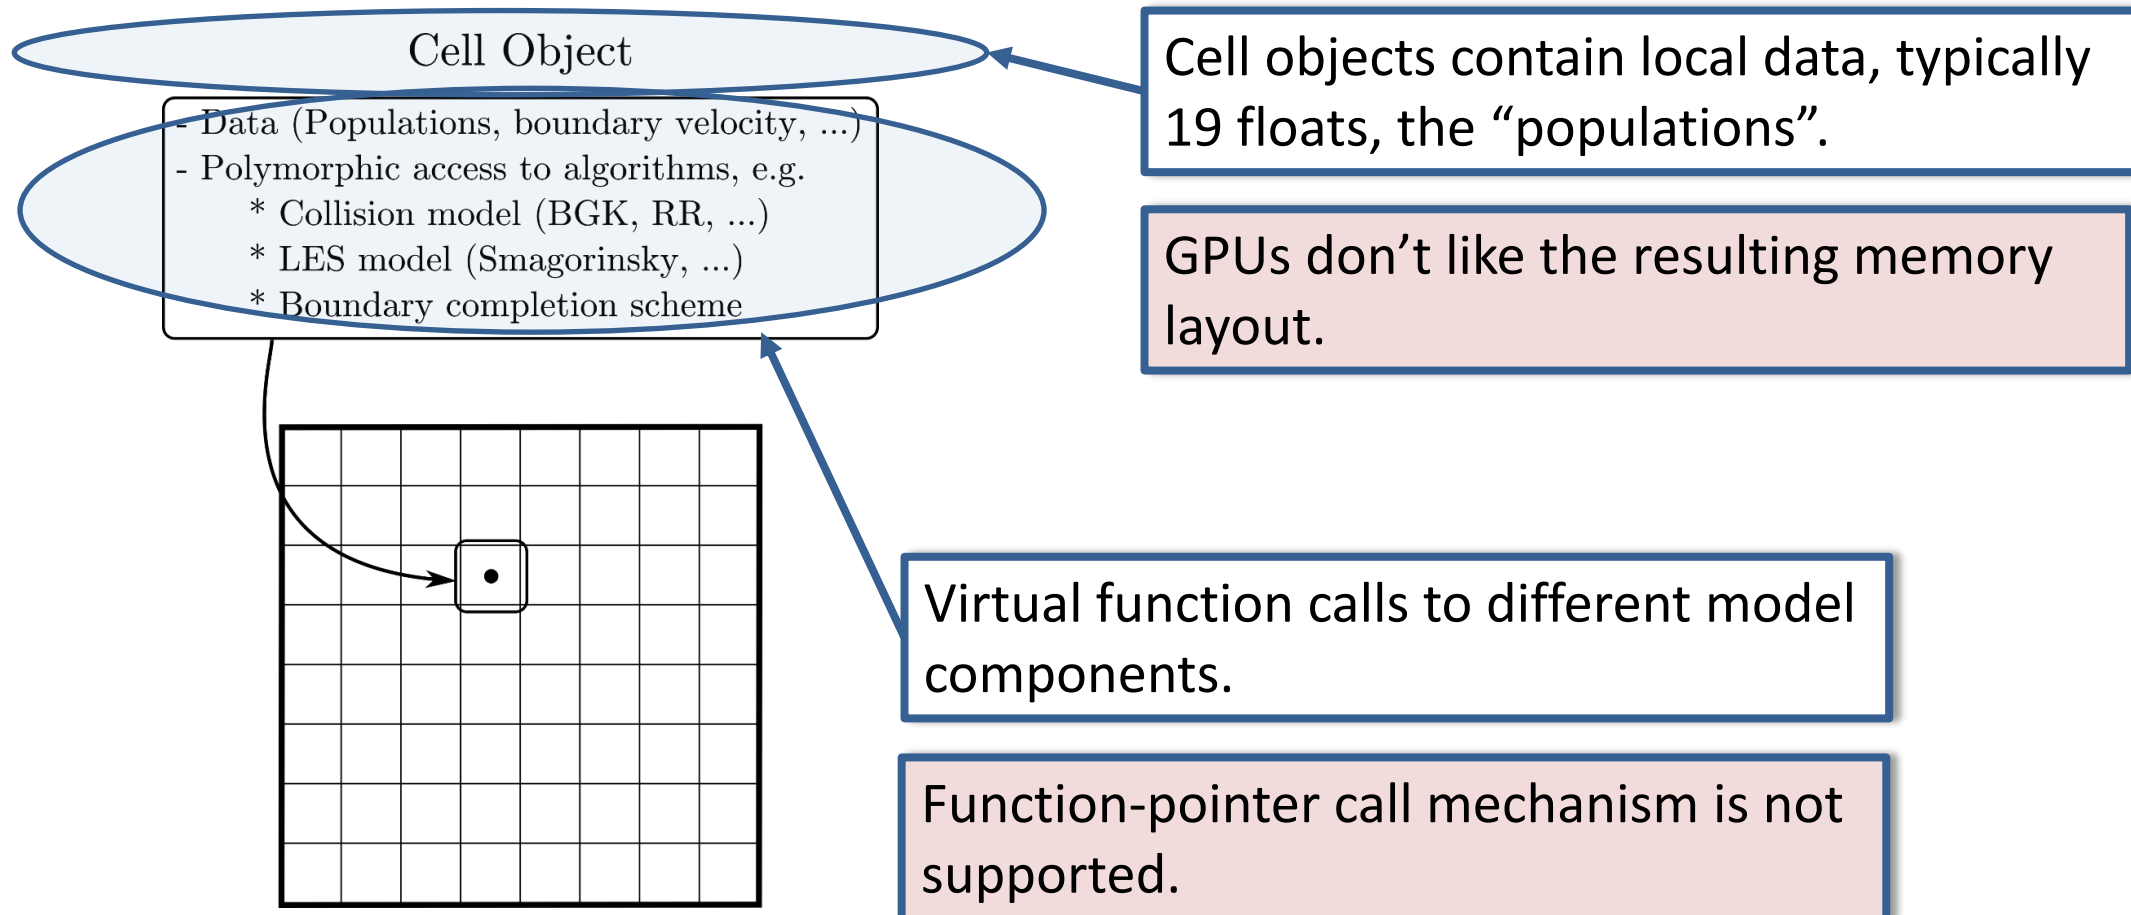


## Part II

# DESIGN: FROM OBJECT-ORIENTED TO DATA-ORIENTED

# Palabos: Object-oriented approach

*Polymorphism allows every grid node to implement different physical / numerical model*

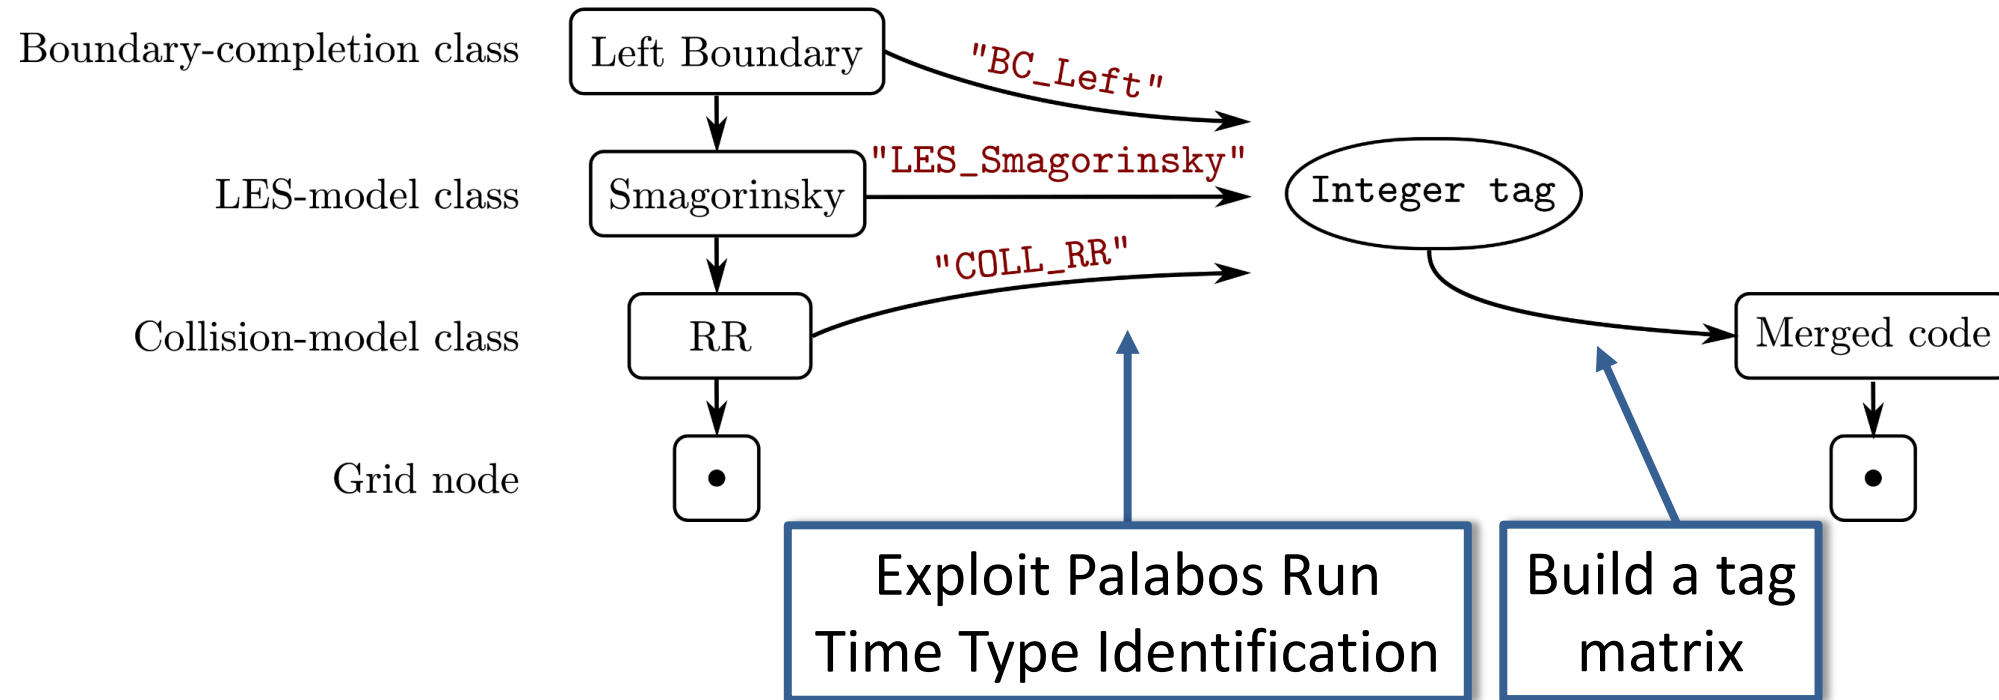


# Polymorphic objects → Tag matrix

*The tag matrix suits the GPU better and can be generated automatically*

On CPU

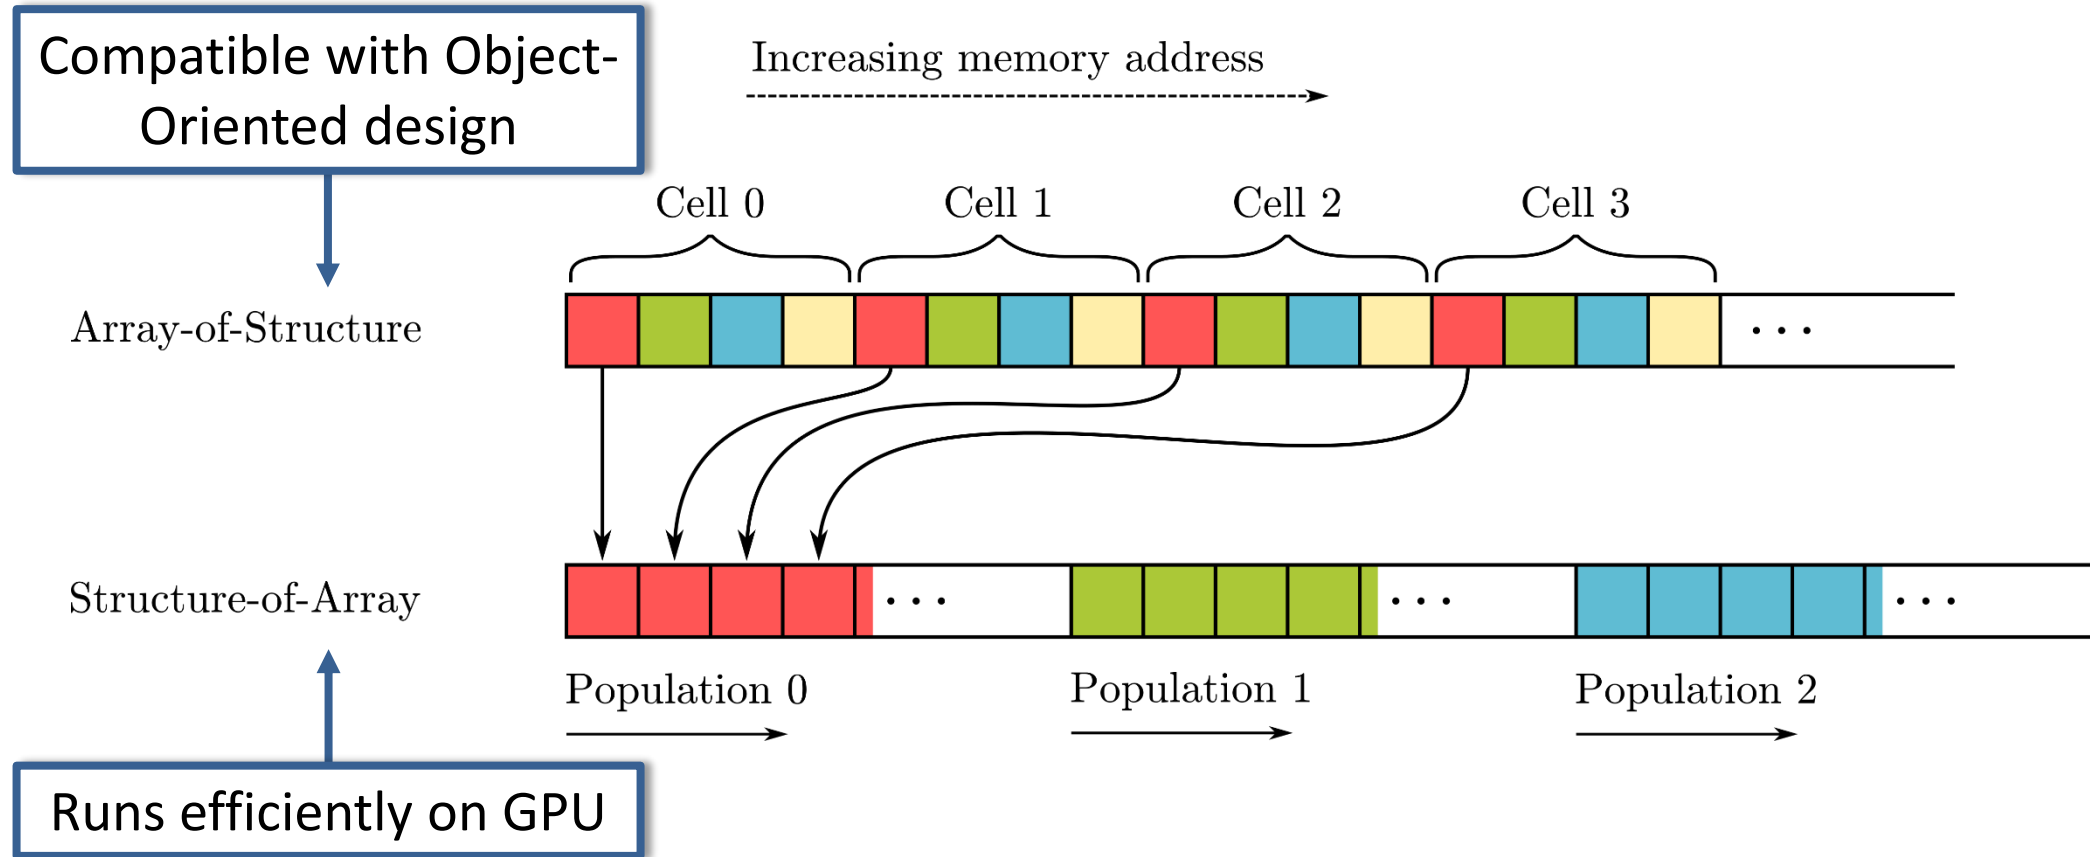
On GPU





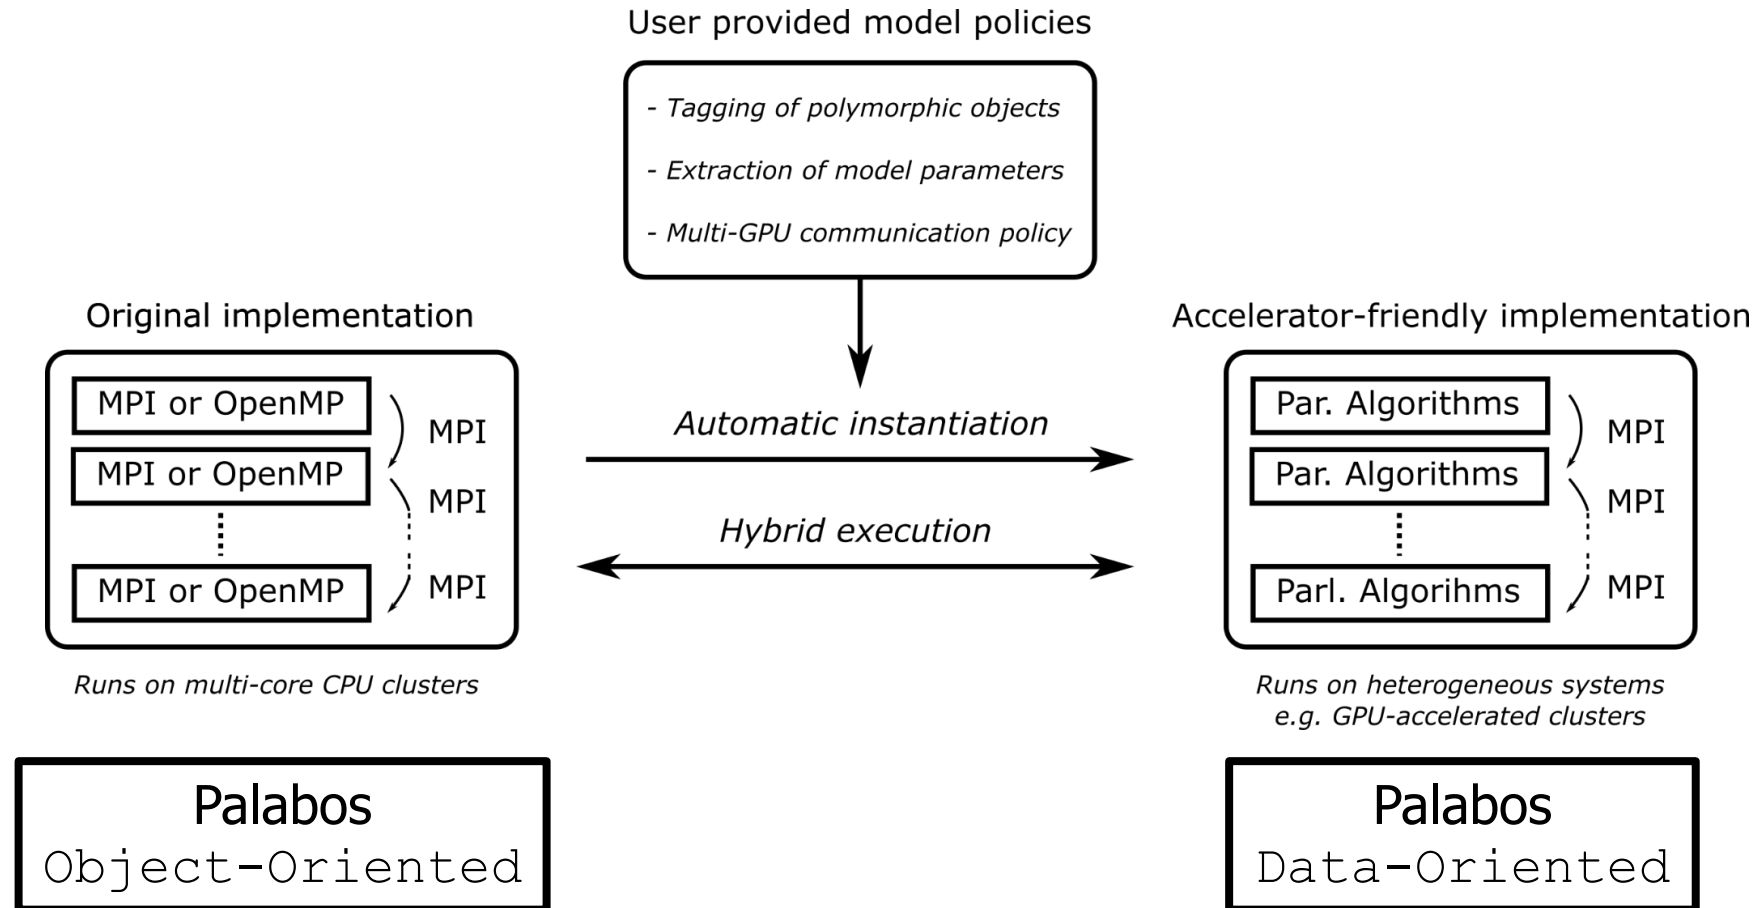
# Array-of-structure → Structure-of-array

*Data alignment in memory must be optimized for the way a node interacts with its neighbors*



# Object-oriented → Data-oriented

*A policy class, provided for each model, allows automatic instantiation of data-oriented structure*



# Benchmark hardware: 2 desktop computers

## **Computer 1**

(Sits on my desk)

Intel Xeon Gold 6240R CPU

48 cores, 2 sockets

## **Computer 2**

(On a desk next door)

NVIDIA RTX 3090 GPU

Ampere Architecture

All benchmarks executed in  
single-precision


# First validation: lid-driven cavity

*With a simple geometry, this standard benchmark assesses the raw performance of the code*

Moving top lid: constant velocity  
from left to right

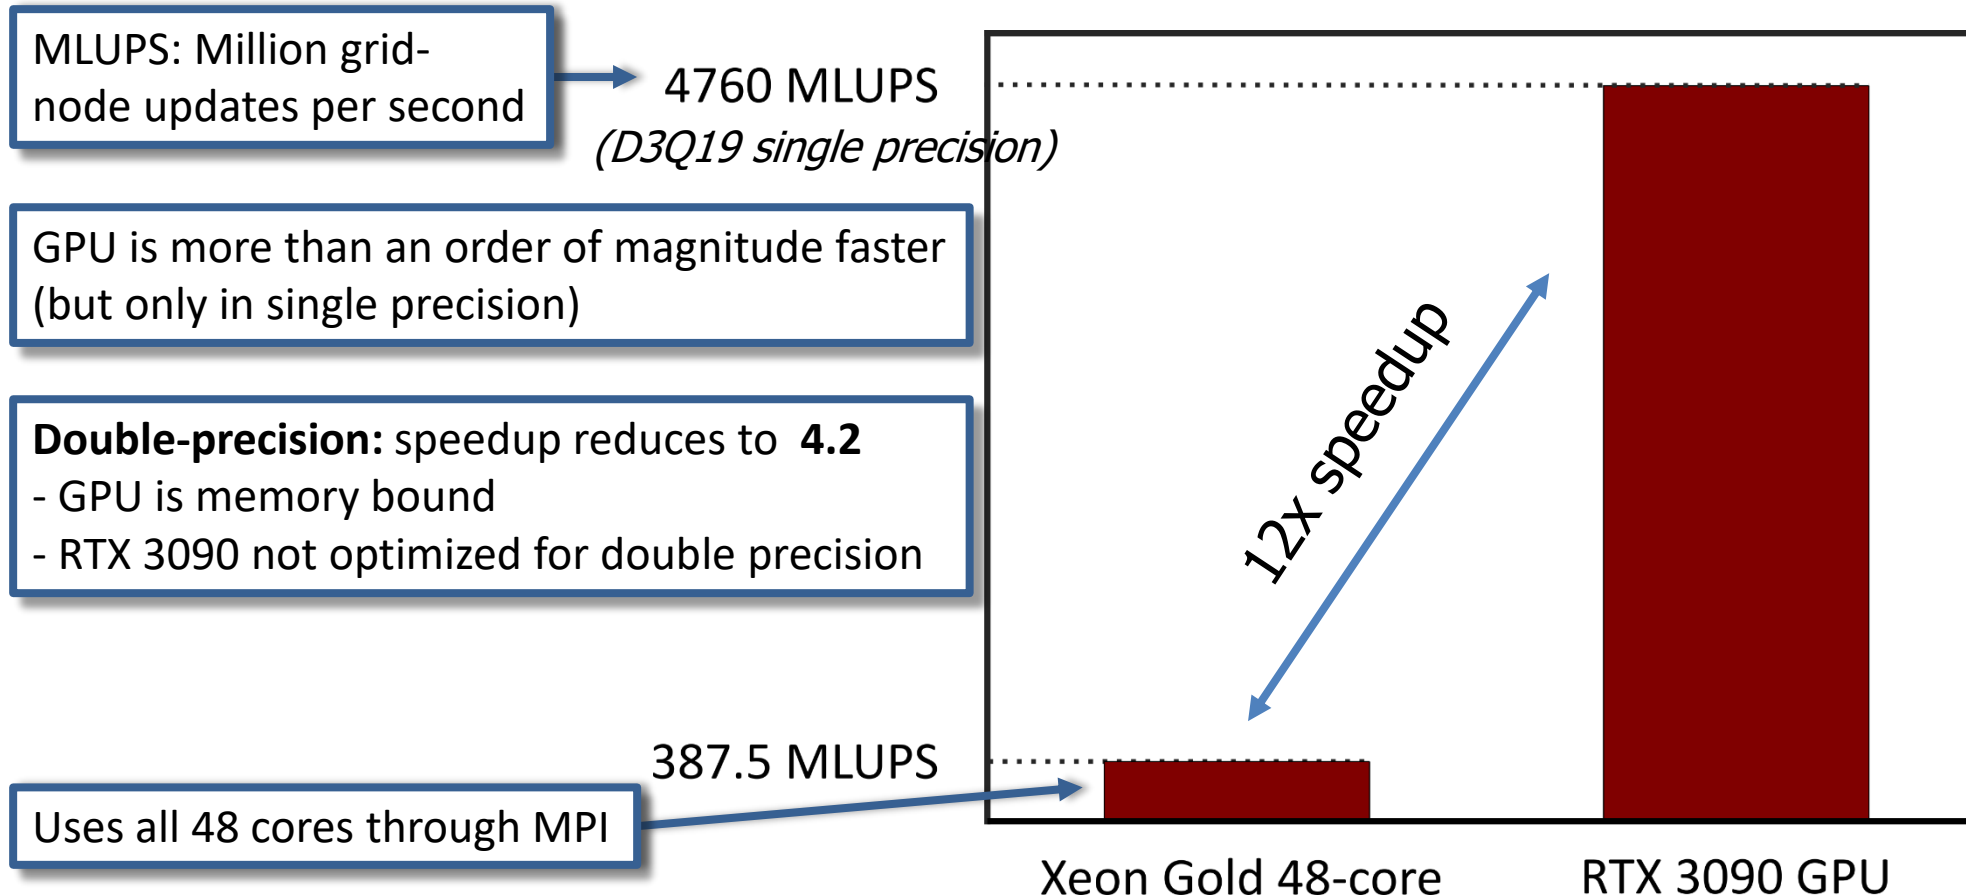
Cubic box with no-slip walls

**Collision:** Recursive-Regularized  
**Lattice:** D3Q19  
**Resolution:** 420x420x420  
**Floats:** Single precision



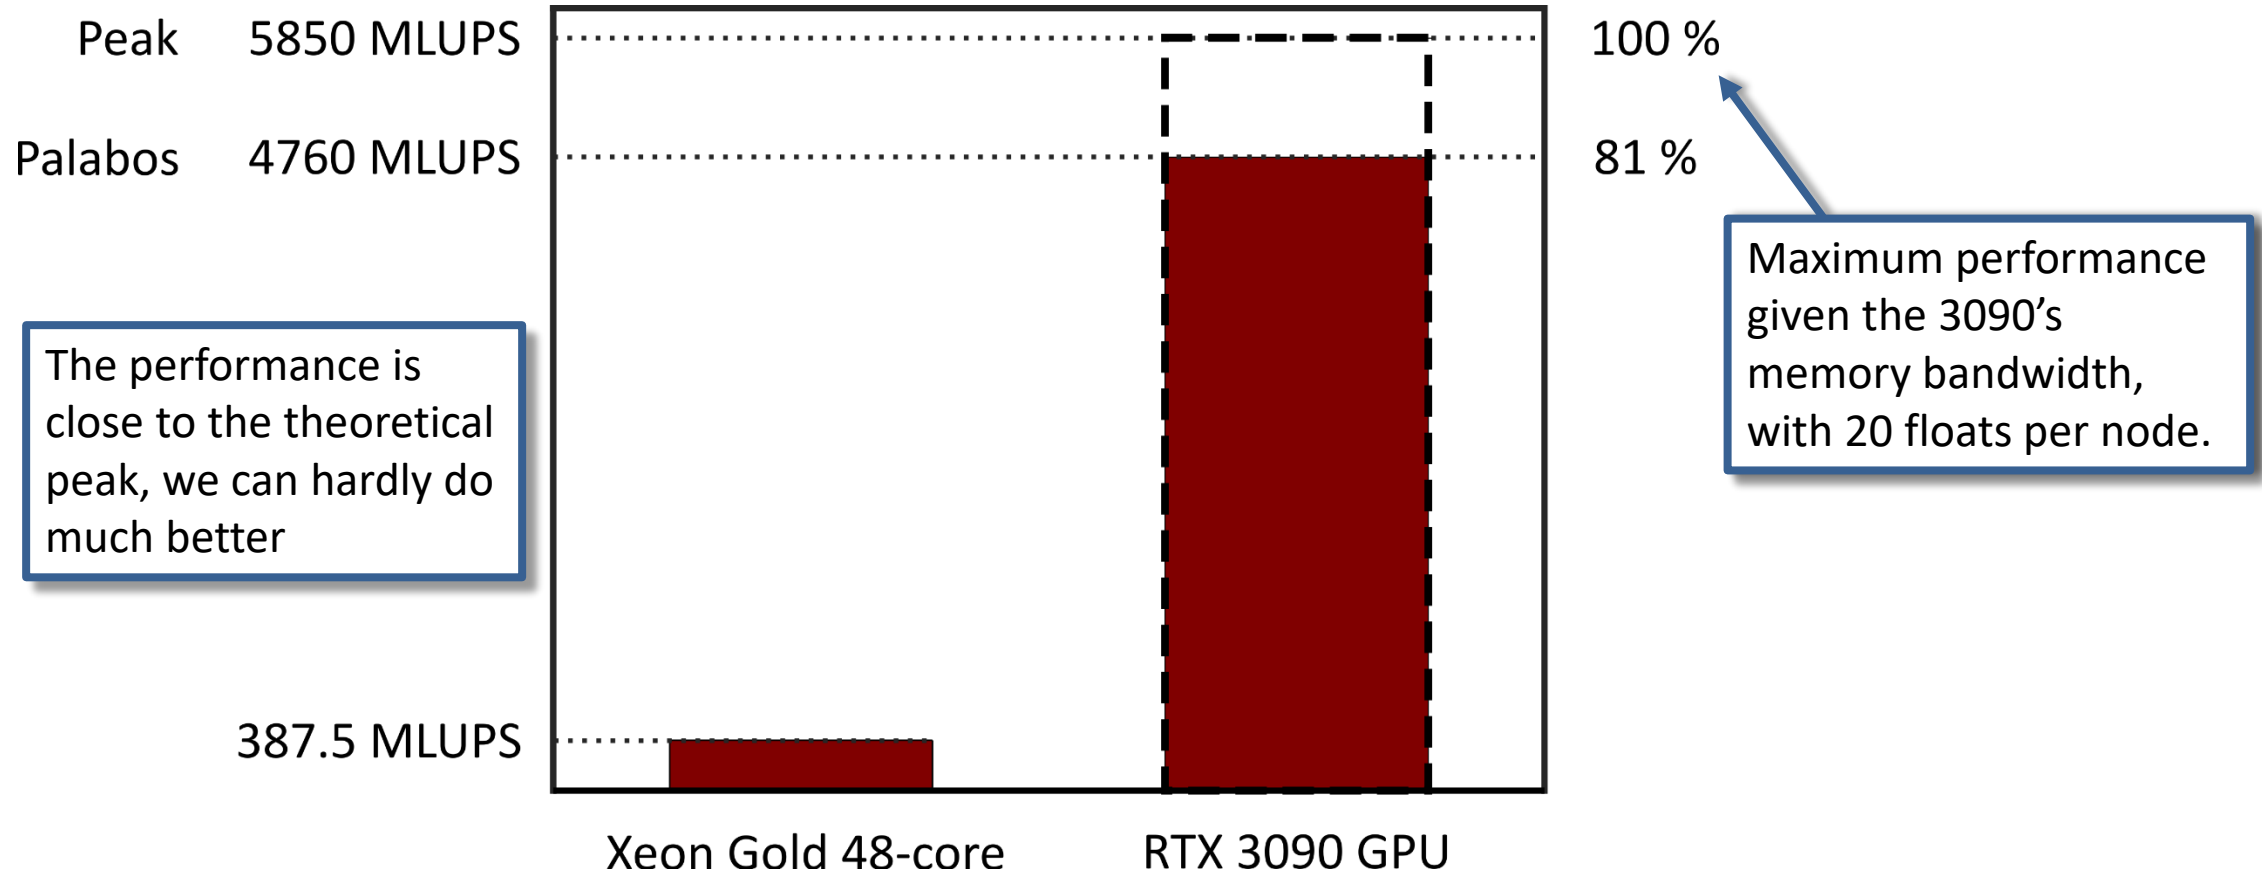
# Performance: lid-driven cavity

*The same end-user application is executed once with the CPU and once with the GPU backend*



# Comparison against peak performance

*Because the problem is memory bound on GPU, peak performance is dictated by memory bandwidth*




# 2<sup>nd</sup> validation: Porous media flow

*A fully resolved flow through a porous rock, requires handling of a very complex geometry*

Handling of geometry: use of different collision model on solid and fluid nodes

In the data-oriented code, how do we store extra data required by specific nodes ?  
Example: flow on inlet condition




# Separate array for extra grid-node data



*Some nodes require extra data, such as the parameters of a boundary condition*

Parameter Array



In object-oriented approach,  
no problem: each node  
freely allocates extra data.

In data-oriented approach,  
extra data is packed in a  
parameter array. Each node  
has an index into this array.



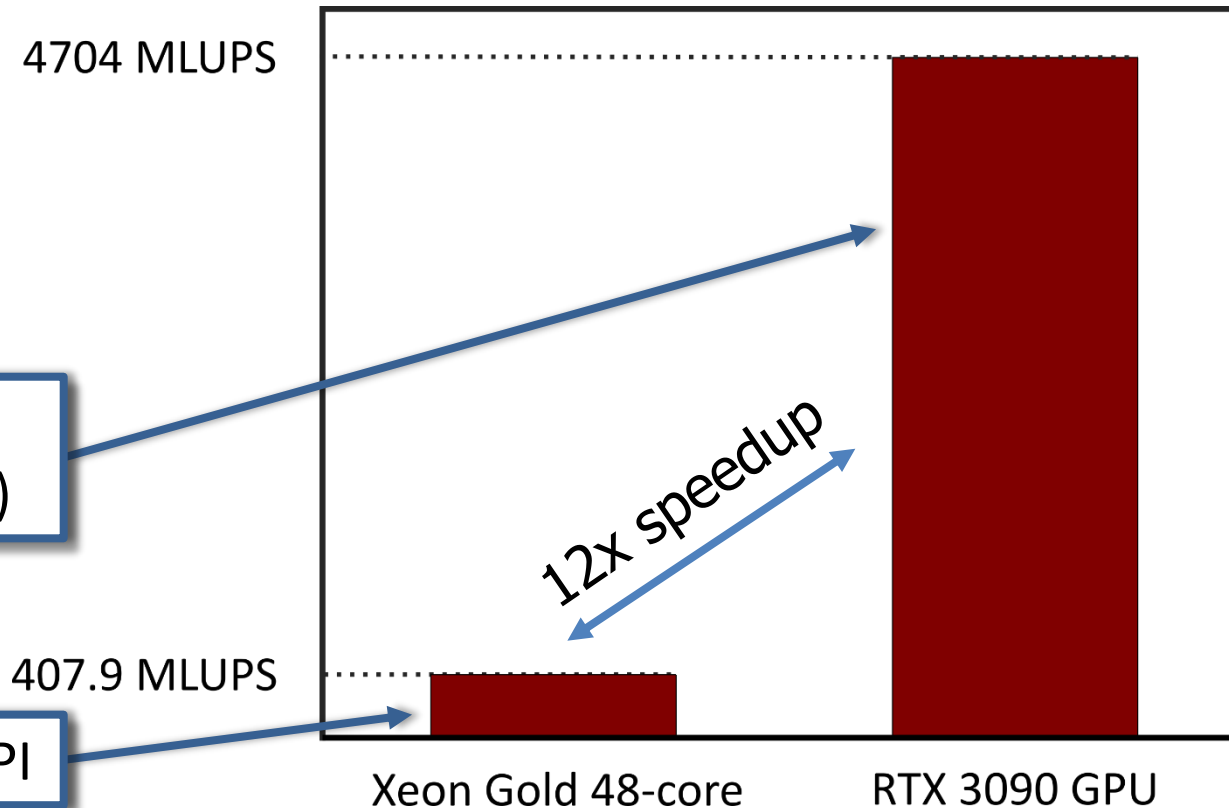
# Performance: Porous media flow

*CPU-to-GPU backend switch yields over 10-fold improvement for this complex example as well*

**Collision:** TRT  
**Lattice:** D3Q19  
**Resolution:** 300x300x340  
**Floats:** Single precision

Performance obtained with  
latest version of nvc++ (21.9)

Uses all 48 cores through MPI




# 3<sup>rd</sup> case: 3D multi-component flow

*The most complex of the 3 examples, this multi-component flow requires coupling between two grids*

An interaction term is evaluated  
to model interface physics

Each fluid component is  
simulated on a separate grid

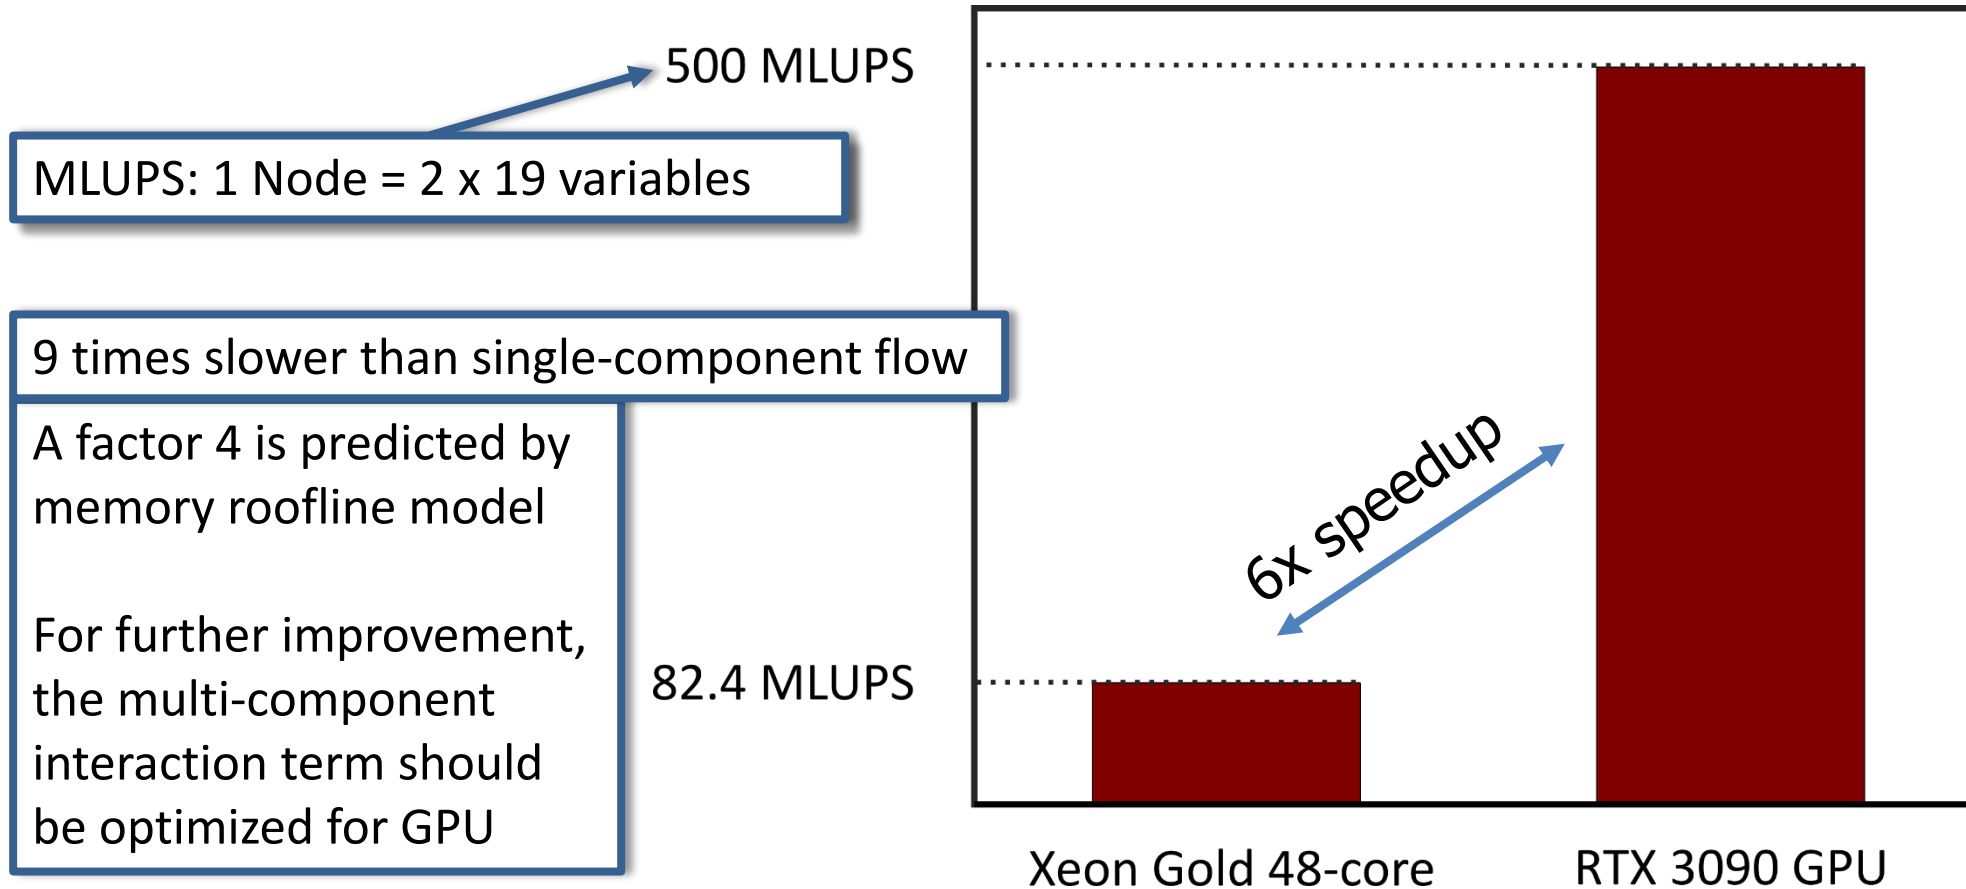


**Collision:** BGK  
**Model:** Pseudo-potential  
**Lattice:** D3Q19  
**Resolution:** 330x330x110  
**Floats:** Single precision

Boundaries: implementation of contact physics (contact angle)

# Performance: 3D multi-component flow

*The problem is stressful for memory-bound hardware; GPU backend still yields substantial speedup*





## Part III

# IN-DEPTH: PERFORMANCE IMPROVEMENTS

# Multi-GPU for a hybrid CPU/GPU code

Some components of Palabos are not ported to GPU yet.

Some components will not be ported to GPU (in particular pre-processing).

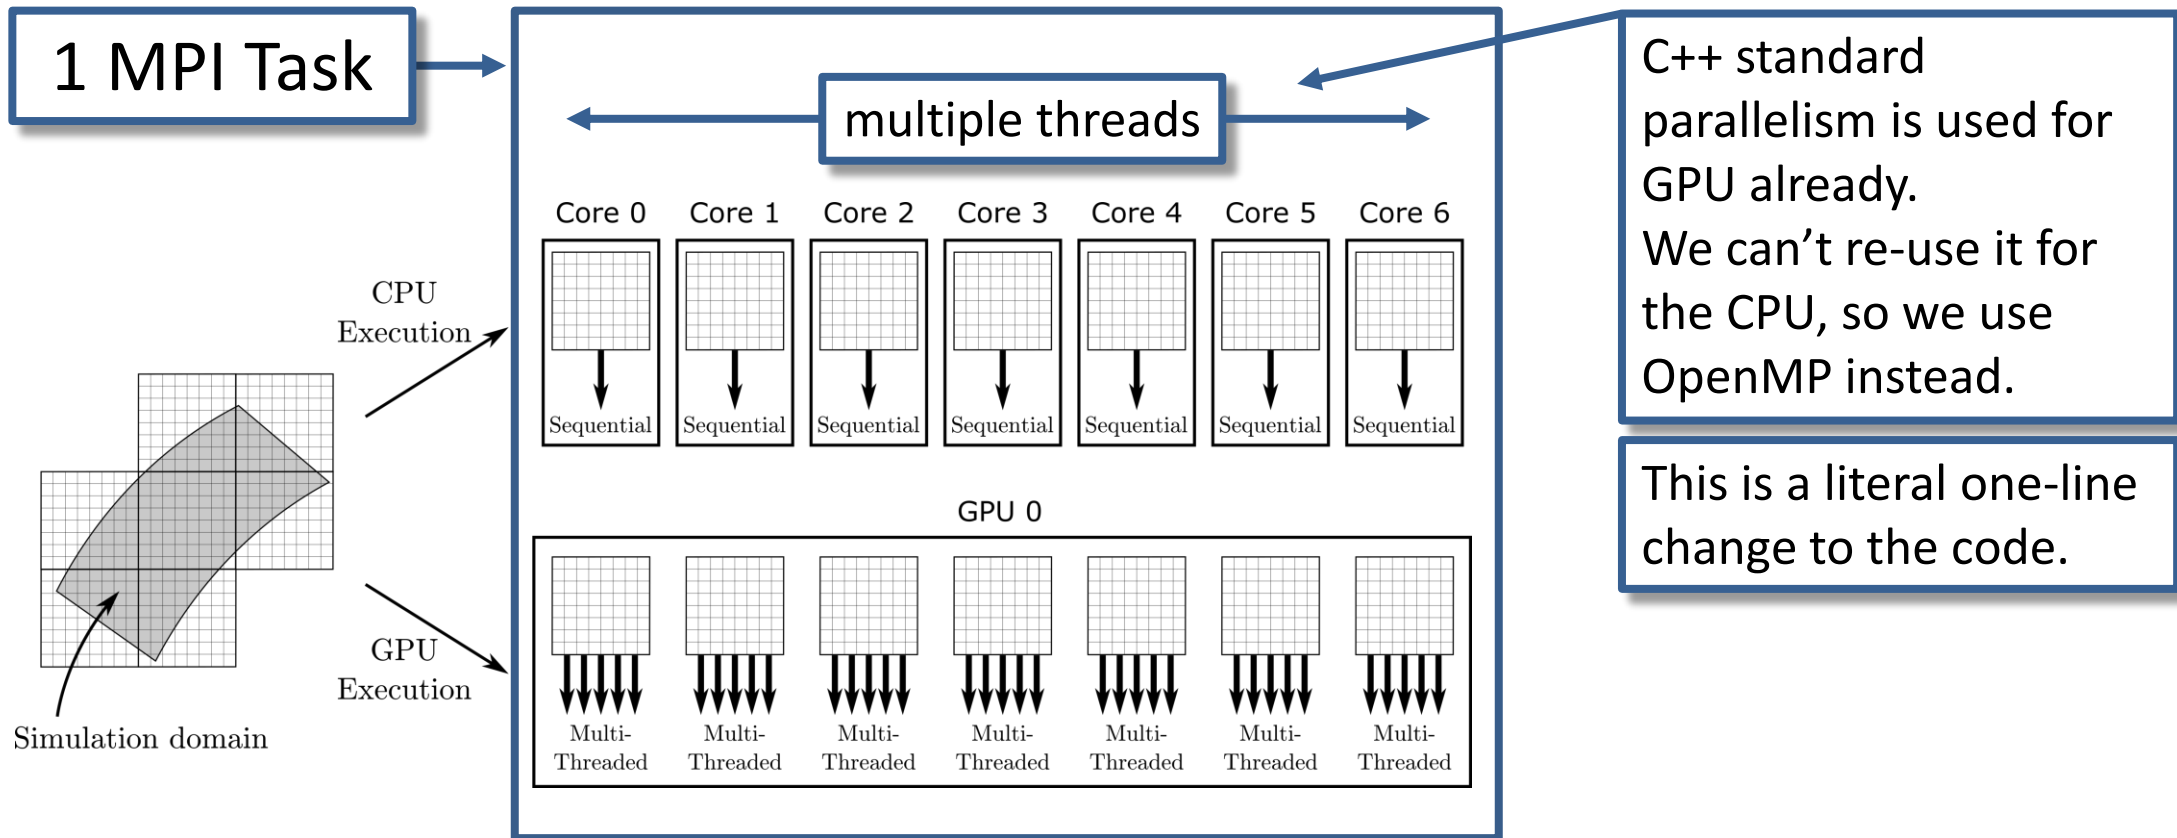
Solution: hybrid CPU / GPU execution.

With standard parallelism: everything is compiled together with nvc++.

For a multi-CPU, multi-GPU code, CPU and GPU are required to have same number of MPI threads. Problem: CPU and GPU have not been parallelized according to the same philosophy.

# Solution: MPI + OpenMP on CPU

*Both CPU and GPU get one MPI task per cluster node, then the CPU gets some OpenMP threads*




# Validation: Multi-GPU execution

*The existing MPI backend can be reused for multi-GPU as is, although performance is not brilliant*

**Lattice Boltzmann Model**  
**Problem:** Lid-driven cavity  
**Collision:** TRT  
**Lattice:** D3Q19  
**Floats:** Single precision

**Cluster: Selene**  
**Nodes:** 8 GPUs  
**GPUs:** A100-SXM4-80GB



**Success:** Multi-GPU works with existing communication layer (packing / unpacking, MPI, ...)

**Issue:** Scaling below state-of-the-art.

**Future improvements:**

- Use of pinned memory for communication buffers.
- Overlap communication and computation.

# Optimization: collision-kernel size

The issue with the tag-matrix approach: somewhere in your code, you find a huge switch statement

```
template<typename T, template<typename U> class Descriptor>
void collide(int collisionModel, Array<T, Descriptor<T>::numPop>& f,
            Array<T, Descriptor<T>::ExternalField::numScalars>& ext,
            Array<T, GPUconst<T, Descriptor>::maxStaticScalars> staticScalars, T* dynamicScalars, plint index)
{
    switch (collisionModel) {
    case CollisionModel::NoDynamics:
        Collision<T, Descriptor, CollisionModel::NoDynamics>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BounceBack:
        Collision<T, Descriptor, CollisionModel::BounceBack>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BGK:
        Collision<T, Descriptor, CollisionModel::BGK>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BGK_ExternalMoment:
        Collision<T, Descriptor, CollisionModel::BGK_ExternalMoment>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::TRT:
        Collision<T, Descriptor, CollisionModel::TRT>::collide(f, ext, staticScalars, dynamicScalars, index);
        // And so on, and so on ...
    }
```



# Optimization: collision-kernel size

*The issue with the tag-matrix approach: somewhere in your code, you find a huge switch statement*

## Problem 1: Maintenance

It is impossible to propose new models without modifying internal Palabos code.

```
template<typename T, template<typename U> class Descriptor>
void collide(int collisionModel, Array<T, Descriptor<T>::numPop>& f,
            Array<T, Descriptor<T>::ExternalField::numScalars>& ext,
            Array<T, GPUconst<T, Descriptor>::maxStaticScalars> staticScalars, T* dynamicScalars, plint index)
{
    switch (collisionModel) {
    case CollisionModel::NoDynamics:
        Collision<T, Descriptor, CollisionModel::NoDynamics>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BounceBack:
        Collision<T, Descriptor, CollisionModel::BounceBack>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BGK:
        Collision<T, Descriptor, CollisionModel::BGK>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::BGK_ExternalMoment:
        Collision<T, Descriptor, CollisionModel::BGK_ExternalMoment>::collide(f, ext, staticScalars, dynamicScalars, index);
        break;
    case CollisionModel::TRT:
        Collision<T, Descriptor, CollisionModel::TRT>::collide(f, ext, staticScalars, dynamicScalars, index);
        // And so on, and so on ...
    }
```

## Problem 2: Size of collision kernels

- Potentially hundreds of collision models.
- Everything gets compiled into a single Cuda kernel.
- Size of Cuda kernels have performance impact.

# Solution: variadic templates

*All models that are actually used are provided to the kernel at compile time using templates*

```
lattice -> collideAndStream (  
    CollisionKernel<T, DESCRIPTOR,  
    CollisionModel::TRT,  
    CollisionModel::BounceBack,  
    CollisionModel::Boundary_RegularizedVelocity_0_1__TRT,  
    CollisionModel::Boundary_RegularizedVelocity_0_M1__TRT>());
```

- All models required in the kernel are provided as template arguments.
- Variadic templates allow a variable number of template arguments.
- “Switch statement” automatically generated in pre-compilation stage.

# Solution: variadic templates

```
template<int MODEL>
static void static_switch(IntList<MODEL>, int collisionModel, Array<T, Descriptor<T>::numPop>& f,
    Array<T, Descriptor<T>::ExternalField::numScalars>& ext, Array<T, GPUconst<T, Descriptor>::maxStaticScalars> staticScalars,
    T* dynamicScalars, plint index)
{
    if (collisionModel == MODEL) {
        Collision<T, Descriptor, MODEL>::collide(f, ext, staticScalars, dynamicScalars, index);
    }
    else {
        printf("Collision model not implemented: %d\n", collisionModel);
    }
}

template<int ...N>
static void static_switch(int collisionModel, Array<T, Descriptor<T>::numPop>& f,
    Array<T, Descriptor<T>::ExternalField::numScalars>& ext, Array<T, GPUconst<T, Descriptor>::maxStaticScalars> staticScalars,
    T* dynamicScalars, plint index)
{
    static_switch(IntList<N...>(), collisionModel, f, ext, staticScalars, dynamicScalars, index);
}
```

Template argument list is processed recursively on every node to match the local tag

Performance gain from using variadic templates, on RTX 3090: 12%

# Conclusion

- C++ standard parallelism allows porting a CPU code to GPU, conveniently, step-by-step, efficiently.
- Ideal for hybrid CPU / GPU codes, including multi-CPU / multi-GPU.
- The approach is entirely portable: support on different types of GPUs to be expected.
- Some low-level implementation difficulties occur, but they are addressed with reasonable efforts.
- Most problems are solved at the level of the software architecture, without technical GPU-specific knowledge.

Get the GPU-enabled branch of Palabos:

<https://palabos.unige.ch/community/cpu-gpu-80-days>